

```

# import libraries

import numpy as np
import pandas as pd

%matplotlib inline
import matplotlib.pyplot as plt
plt.style.use('fivethirtyeight')

from sklearn.linear_model import LinearRegression, Lasso, LassoCV,
Ridge, RidgeCV
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, KFold
from sklearn.metrics import r2_score, mean_squared_error,
confusion_matrix, ConfusionMatrixDisplay, accuracy_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import plot_tree

from scipy.optimize import minimize
from scipy.cluster.hierarchy import dendrogram, linkage, fcluster

from tqdm import notebook

import seaborn as sns

```

LOAD DATAFRAME and INITIAL EXPLORATION

```

# loads dataframe; prints head
pisa_df =
pd.read_csv("https://raw.githubusercontent.com/babnigg/DATA11900/main/
economics_and_education_dataset_CSV.csv")
pisa_df.head()

```

	index_code	expenditure_on_education_pct_gdp	mortality_rate_infant
0	AUS-2003	5.246357	4.9
1	AUS-2003	5.246357	4.9
2	AUS-2003	5.246357	4.9
3	AUS-2006	4.738430	4.7
4	AUS-2006	4.738430	4.7

	gini_index	gdp_per_capita_ppp	inflation_consumer_prices
\			

0	33.5	30121.818418	2.732596
1	33.5	30121.818418	2.732596
2	33.5	30121.818418	2.732596
3	NaN	34846.715630	3.555288
4	NaN	34846.715630	3.555288

	intentional_homicides	unemployment	gross_fixed_capital_formation
\			
0	1.533073	5.933	26.050295
1	1.533073	5.933	26.050295
2	1.533073	5.933	26.050295
3	1.372940	4.785	27.789132
4	1.372940	4.785	27.789132

	population_density	suicide_mortality_rate	tax_revenue	\
0	2.567036	10.5	24.299970	
1	2.567036	10.5	24.299970	
2	2.567036	10.5	24.299970	
3	2.662089	10.6	24.511772	
4	2.662089	10.6	24.511772	

	taxes_on_income_profits_capital	alcohol_consumption_per_capita	\
0	62.726546	NaN	
1	62.726546	NaN	
2	62.726546	NaN	
3	65.231562	NaN	
4	65.231562	NaN	

	government_health_expenditure_pct_gdp	urban_population_pct_total
country \		
0	5.623778	84.343
AUS		
1	5.623778	84.343
AUS		
2	5.623778	84.343
AUS		
3	5.719998	84.700
AUS		
4	5.719998	84.700
AUS		

	time	sex	rating
0	2003	BOY	527.0
1	2003	GIRL	522.0
2	2003	TOT	524.0

```
3 2006 BOY 527.0
4 2006 GIRL 513.0
```

```
# filters out dataframe to only include total sex aggregate; prints info and describes numerical variables
```

```
agg = "TOT"
```

```
pisa_df_tot = pisa_df[pisa_df["sex"]==agg]
```

```
pisa_df_tot.info()
```

```
pisa_df_tot.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 214 entries, 2 to 633
```

```
Data columns (total 20 columns):
```

#	Column	Non-Null Count	Dtype
0	index_code	214 non-null	object
1	expenditure_on_education_pct_gdp	202 non-null	float64
2	mortality_rate_infant	214 non-null	float64
3	gini_index	181 non-null	float64
4	gdp_per_capita_ppp	214 non-null	float64
5	inflation_consumer_prices	214 non-null	float64
6	intentional_homicides	187 non-null	float64
7	unemployment	214 non-null	float64
8	gross_fixed_capital_formation	214 non-null	float64
9	population_density	214 non-null	float64
10	suicide_mortality_rate	214 non-null	float64
11	tax_revenue	198 non-null	float64
12	taxes_on_income_profits_capital	198 non-null	float64
13	alcohol_consumption_per_capita	37 non-null	float64
14	government_health_expenditure_pct_gdp	214 non-null	float64
15	urban_population_pct_total	214 non-null	float64
16	country	214 non-null	object
17	time	214 non-null	int64
18	sex	214 non-null	object
19	rating	214 non-null	float64

```
dtypes: float64(16), int64(1), object(3)
```

```
memory usage: 35.1+ KB
```

	expenditure_on_education_pct_gdp	mortality_rate_infant
gini_index \		
count	202.000000	214.000000
181.000000		
mean	5.237384	5.062150
33.411050		
std	1.105318	4.090917
6.816531		
min	3.040150	1.900000
24.400000		
25%	4.460890	3.000000

28.100000		
50%	5.058085	3.700000
32.700000		
75%	5.796320	5.100000
35.400000		
max	8.448880	25.300000
57.600000		

	gdp_per_capita_ppp	inflation_consumer_prices
intentional_homicides \		
count	214.000000	214.000000
187.000000		
mean	36176.526959	2.300823
3.119124		
std	16451.684140	2.742357
6.122298		
min	9587.557951	-4.478103
0.000000		
25%	26081.142138	0.799862
0.833583		
50%	34305.514059	1.980056
1.095505		
75%	42993.833915	2.991248
1.842371		
max	116498.512081	21.602438
29.581371		

	unemployment	gross_fixed_capital_formation	population_density
\			
count	214.000000	214.000000	214.000000
mean	7.643659	22.507161	133.327098
std	4.021760	3.942642	133.840583
min	2.246000	10.770040	2.567036
25%	4.892500	20.210197	30.425942
50%	6.843000	22.454439	101.872663
75%	9.333500	24.296864	192.362949
max	24.981000	36.800286	528.969011

	suicide_mortality_rate	tax_revenue
taxes_on_income_profits_capital \		
count	214.000000	198.000000
198.000000		


mean	13.584112	20.119590
27.827355		
std	6.282140	5.748563
13.303791		
min	2.200000	7.903518
7.173176		
25%	9.675000	15.526577
17.681421		
50%	12.900000	20.951774
25.383772		
75%	16.925000	24.172369
34.128627		
max	34.200000	33.921619
65.517310		

	alcohol_consumption_per_capita	government_health_expenditure_pct_gdp \
count	214.000000	37.000000
mean	9.452505	
5.939307		
std	2.721759	
1.755618		
min	1.895130	
2.413703		
25%	7.847910	
4.473900		
50%	9.916870	
6.045355		
75%	11.462780	
7.382606		
max	15.058720	
9.276339		

	urban_population_pct_total	time	rating
count	214.000000	214.000000	214.000000
mean	77.345145	2010.808411	489.242991
std	11.020002	5.106906	37.238186
min	51.758000	2003.000000	356.000000
25%	68.733000	2006.000000	482.000000
50%	79.628500	2012.000000	495.000000
75%	85.680750	2015.000000	512.750000
max	98.001000	2018.000000	554.000000

```
# plots histogram of PISA score distribution in original data
plt.figure().set_facecolor("lightgrey")
plt.hist(pisa_df_tot.rating, color = "orange")
plt.title("PISA Score Distribution (Original Data)")
plt.ylabel("Number of Observations")
```

```
plt.xlabel("PISA Score")
plt.show()
```



DATA CLEANING and EXPLORATION

```
# counts and prints number of NA values per column
np.sum(pisa_df_tot.isna(),axis=0)
```

index_code	0
expenditure_on_education_pct_gdp	12
mortality_rate_infant	0
gini_index	33
gdp_per_capita_ppp	0
inflation_consumer_prices	0
intentional_homicides	27
unemployment	0
gross_fixed_capital_formation	0

```

population_density          0
suicide_mortality_rate      0
tax_revenue                 16
taxes_on_income_profits_capital 16
alcohol_consumption_per_capita 177
government_health_expenditure_pct_gdp 0
urban_population_pct_total  0
country                     0
time                        0
sex                         0
rating                      0
dtype: int64

```

```

# defining function to impute data with the mean for the column
def fill_missing_with_mean(group):
    '''fill the missing values with the mean of each feature (column)
    for each country.'''
    columns_to_fill = ['expenditure_on_education_pct_gdp',
                       'gini_index', 'intentional_homicides',
                       'tax_revenue',
                       'taxes_on_income_profits_capital', 'alcohol_consumption_per_capita']
    for column in columns_to_fill:
        group[column] = group[column].fillna(group[column].mean())
    return group

# apply the function to each group (country)
pisa_df_tot = pisa_df_tot.groupby('country',
group_keys=False).apply(fill_missing_with_mean)

# glance at data head after imputing
display(pisa_df_tot.head())

# counts and prints number of NA values per column after imputing
np.sum(pisa_df_tot.isna(),axis=0)

```

```

index_code  expenditure_on_education_pct_gdp
mortality_rate_infant \
2    AUS-2003                      5.246357
4.9
5    AUS-2006                      4.738430
4.7
8    AUS-2009                      5.081320
4.2
11   AUS-2012                      4.866670
3.6
14   AUS-2015                      5.315520
3.3

```

```

2    gini_index  gdp_per_capita_ppp  inflation_consumer_prices \
2          33.5      30121.818418          2.732596

```

5	33.9	34846.715630	3.555288
8	33.9	40312.395119	1.771117
11	33.9	42866.604330	1.762780
14	33.9	46292.095439	1.508367
intentional_homicides unemployment gross_fixed_capital_formation			
\			
2	1.533073	5.933	26.050295
5	1.372940	4.785	27.789132
8	1.214170	5.565	27.601846
11	1.069106	5.225	27.404168
14	0.990754	6.055	26.202336
population_density suicide_mortality_rate tax_revenue \			
2	2.567036	10.5	24.299970
5	2.662089	10.6	24.511772
8	2.823588	11.2	22.021006
11	2.959200	11.7	21.097483
14	3.100113	13.2	21.866424
taxes_on_income_profits_capital alcohol_consumption_per_capita \			
2	62.726546	10.3903	
5	65.231562	10.3903	
8	64.951075	10.3903	
11	65.517310	10.3903	
14	64.893433	10.3903	
government_health_expenditure_pct_gdp urban_population_pct_total			
country \			
2	5.623778	84.343	
AUS			
5	5.719998	84.700	
AUS			
8	6.244110	85.063	
AUS			
11	6.179476	85.402	
AUS			
14	7.234067	85.701	
AUS			
time sex rating			
2	2003	TOT	524.0
5	2006	TOT	520.0
8	2009	TOT	514.0


```

11  2012  TOT    504.0
14  2015  TOT    494.0

index_code                                0
expenditure_on_education_pct_gdp          0
mortality_rate_infant                     0
gini_index                                12
gdp_per_capita_ppp                        0
inflation_consumer_prices                 0
intentional_homicides                     6
unemployment                              0
gross_fixed_capital_formation             0
population_density                        0
suicide_mortality_rate                    0
tax_revenue                               6
taxes_on_income_profits_capital           6
alcohol_consumption_per_capita            2
government_health_expenditure_pct_gdp     0
urban_population_pct_total                 0
country                                   0
time                                       0
sex                                         0
rating                                     0
dtype: int64

```

Since there are some countries that don't have any entry for certain variables we can just drop those countries.

Countries: BEL, JPN, NZL, CRI, LTU.

We also choose to drop countries with the characteristically low PISA scores in comparison to other countries in dataset to normalize the score distribution.

Countries: MEX, BRA, CHL, TUR, COL

```

# dropping countries with lone NaN values
# last 5 countries in list are lowest scoring countries
countries_to_drop = ['BEL', 'JPN',
                    'NZL', 'CRI', 'LTU', 'MEX', 'BRA', 'CHL', 'TUR', 'COL']
pisa_df_tot =
pisa_df_tot.drop(pisa_df_tot[pisa_df_tot['country'].isin(countries_to_
drop)].index)

# 'sex' is a redundant column as we already control the dataset on
sex='TOT'
pisa_df_tot = pisa_df_tot.drop(['sex'],axis=1)

# informative 'index_code' column can be used to replace default index
pisa_df_tot.reset_index(drop=True,inplace=True)
pisa_df_tot.set_index('index_code',inplace=True)

# printing remaining number of countries in the dataset

```

```
print("Number of remaining countries in
dataset:", len(pisa_df_tot.country.unique()))
print("\n")
```

```
# glance at data head after cleaning
display(pisa_df_tot.head())
```

```
# counts and prints number of NA values per column after removing
countries
print(np.sum(pisa_df_tot.isna(), axis=0))
```

Number of remaining countries in dataset: 29

	expenditure_on _education_pct_gdp	
mortality_rate_infant \		
index_code		
AUS-2003	5.246357	4.9
AUS-2006	4.738430	4.7
AUS-2009	5.081320	4.2
AUS-2012	4.866670	3.6
AUS-2015	5.315520	3.3

	gini_index	gdp_per_capita_ppp	inflation_consumer_prices
\			
index_code			
AUS-2003	33.5	30121.818418	2.732596
AUS-2006	33.9	34846.715630	3.555288
AUS-2009	33.9	40312.395119	1.771117
AUS-2012	33.9	42866.604330	1.762780
AUS-2015	33.9	46292.095439	1.508367

	intentional_homicides	unemployment	\
index_code			
AUS-2003	1.533073	5.933	
AUS-2006	1.372940	4.785	
AUS-2009	1.214170	5.565	
AUS-2012	1.069106	5.225	

AUS-2015	0.990754	6.055
gross_fixed_capital_formation population_density \		
index_code		
AUS-2003	26.050295	2.567036
AUS-2006	27.789132	2.662089
AUS-2009	27.601846	2.823588
AUS-2012	27.404168	2.959200
AUS-2015	26.202336	3.100113

suicide_mortality_rate tax_revenue \		
index_code		
AUS-2003	10.5	24.299970
AUS-2006	10.6	24.511772
AUS-2009	11.2	22.021006
AUS-2012	11.7	21.097483
AUS-2015	13.2	21.866424

taxes_on_income_profits_capital	
alcohol_consumption_per_capita \	
index_code	
AUS-2003	62.726546
10.3903	
AUS-2006	65.231562
10.3903	
AUS-2009	64.951075
10.3903	
AUS-2012	65.517310
10.3903	
AUS-2015	64.893433
10.3903	

government_health_expenditure_pct_gdp	
urban_population_pct_total \	
index_code	
AUS-2003	5.623778
84.343	
AUS-2006	5.719998
84.700	
AUS-2009	6.244110
85.063	
AUS-2012	6.179476
85.402	
AUS-2015	7.234067
85.701	

country time rating	
index_code	

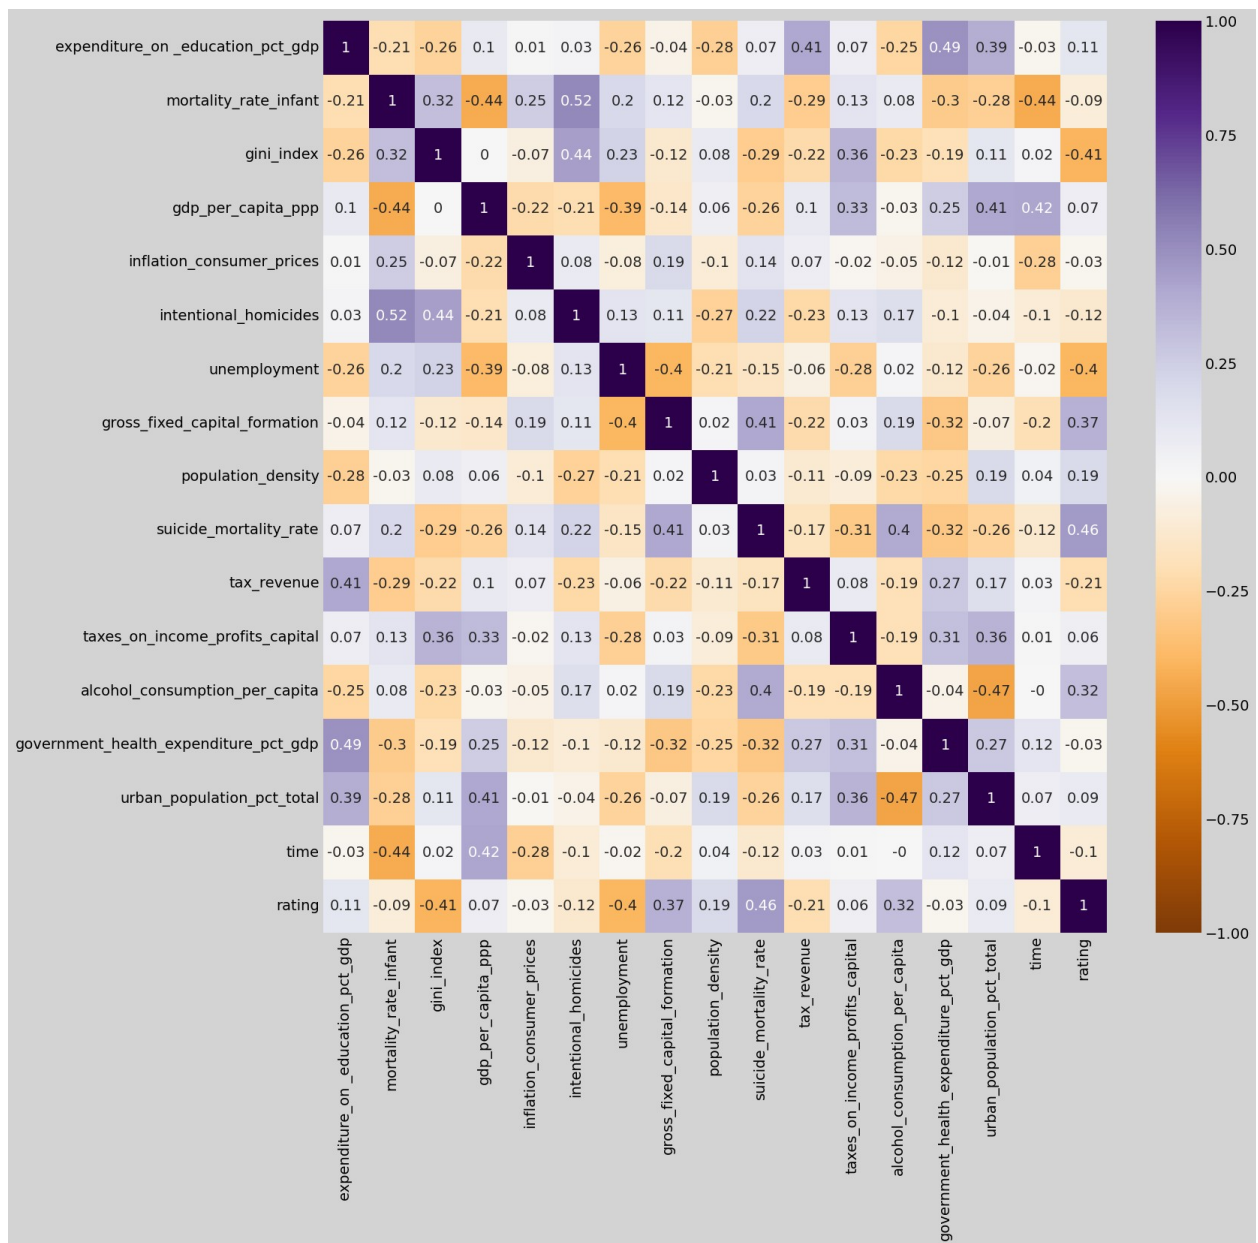
AUS-2003	AUS	2003	524.0
AUS-2006	AUS	2006	520.0
AUS-2009	AUS	2009	514.0
AUS-2012	AUS	2012	504.0
AUS-2015	AUS	2015	494.0

expenditure_on_education_pct_gdp	0
mortality_rate_infant	0
gini_index	0
gdp_per_capita_ppp	0
inflation_consumer_prices	0
intentional_homicides	0
unemployment	0
gross_fixed_capital_formation	0
population_density	0
suicide_mortality_rate	0
tax_revenue	0
taxes_on_income_profits_capital	0
alcohol_consumption_per_capita	0
government_health_expenditure_pct_gdp	0
urban_population_pct_total	0
country	0
time	0
rating	0

dtype: int64

EXPLORATION

```
# plot correlation matrix with all numerical variables
corr_matrix = pisa_df_tot.corr(numeric_only=True).round(2)
plt.figure(figsize=(15,15)).set_facecolor("lightgrey")
sns.heatmap(data=corr_matrix,annot=True,vmin=-1,vmax=1,cmap="PuOr")
plt.show()
```



```
# exploratory data analysis based on correlation matrix


# list of features that has quite high correlation score
features = ["mortality_rate_infant", "gini_index",
            "intentional_homicides",
            "inflation_consumer_prices",
            "government_health_expenditure_pct_gdp", "intentional_homicides",
            "alcohol_consumption_per_capita", "unemployment",
            "suicide_mortality_rate"
          ]
colors = ['blue', 'orange', 'green', 'red', 'purple', 'brown',
          'black', 'gray', 'olive']
```

```
# create subplots
fig, axes = plt.subplots(3, 3, figsize=(18, 18))

# flatten axes for easy iteration
axes = axes.flatten()

# plot scatter plots for each feature
for i, (feature, color) in enumerate(zip(features, colors)):
    ax = axes[i]
    ax.scatter(pisa_df_tot[feature],
pisa_df_tot['rating'], color=color)
    ax.set_xlabel(feature)
    ax.set_ylabel('rating')
    ax.set_title(f'{feature} vs rating')

# adjust layout
plt.tight_layout()
plt.show()
```



```
# scaling and transforming numerical variables
```

```
scaler = StandardScaler()
```

```
num_X = pisa_df_tot.select_dtypes(exclude='object')
```

```
num_columns = num_X.drop(['rating', 'time'], axis=1).columns
```

```
# (time is a number, but it is indeed a categorical variable!)
```

```
# uses standard scaler to transform all numerical variables
```

```
pisa_df_tot[num_columns]=
```

```
scaler.fit_transform(pisa_df_tot[num_columns])
```

```
# one-hot-encodes categorical variables of country and time
```

```
country_col = pisa_df_tot.country
```

```
pisa_df_tot =
pd.get_dummies(pisa_df_tot,columns=['country','time'],drop_first=True,
dtype=int)
pisa_df_tot["country"] = country_col
display(pisa_df_tot.head())
```

	expenditure_on _education_pct_gdp	mortality_rate_infant \
index_code		
AUS-2003	-0.029140	0.831882
AUS-2006	-0.496382	0.679648
AUS-2009	-0.180958	0.299063
AUS-2012	-0.378414	-0.157639
AUS-2015	0.034482	-0.385990

	gini_index	gdp_per_capita_ppp	inflation_consumer_prices
\			
index_code			
AUS-2003	0.462126	-0.546819	0.507257
AUS-2006	0.559458	-0.258194	0.968513
AUS-2009	0.559458	0.075682	-0.031812
AUS-2012	0.559458	0.231708	-0.036486
AUS-2015	0.559458	0.440957	-0.179127

	intentional_homicides	unemployment	\
index_code			
AUS-2003	0.142289	-0.426994	
AUS-2006	-0.005489	-0.695237	
AUS-2009	-0.152009	-0.512982	
AUS-2012	-0.285880	-0.592426	
AUS-2015	-0.358186	-0.398488	

	gross_fixed_capital_formation	population_density	\
index_code			
AUS-2003	0.872513	-1.001173	
AUS-2006	1.294253	-1.000452	
AUS-2009	1.248829	-0.999228	
AUS-2012	1.200883	-0.998200	
AUS-2015	0.909389	-0.997132	


	suicide_mortality_rate	...	country_SVK	country_SVN	\
index_code		...			
AUS-2003	-0.660127	...	0	0	
AUS-2006	-0.642575	...	0	0	
AUS-2009	-0.537262	...	0	0	
AUS-2012	-0.449501	...	0	0	
AUS-2015	-0.186219	...	0	0	

	country_SWE	country_USA	time_2006	time_2009	time_2012
\					
index_code					
AUS-2003	0	0	0	0	0
AUS-2006	0	0	1	0	0
AUS-2009	0	0	0	1	0
AUS-2012	0	0	0	0	1
AUS-2015	0	0	0	0	0

	time_2015	time_2018	country
index_code			
AUS-2003	0	0	AUS
AUS-2006	0	0	AUS
AUS-2009	0	0	AUS
AUS-2012	0	0	AUS
AUS-2015	1	0	AUS

[5 rows x 50 columns]

```
# plots histogram of PISA score distribution after data cleaning
plt.figure().set_facecolor("lightgrey")
plt.hist(pisa_df_tot.rating, color = "orange")
plt.title("PISA Score Distribution")
plt.ylabel("Number of Observations")
plt.xlabel("PISA Score")
plt.show()
```



LINEAR, RIDGE, and LASSO MODELS: SELECTING BEST MODEL

```
# train and test dataset split
x_df = pisa_df_tot.loc[:,pisa_df_tot.columns != "rating"]
y_df = pisa_df_tot["rating"]

x_train, x_test, y_train, y_test = train_test_split(x_df, y_df,
test_size = 0.2)

# defining 7 different models
model1x = x_train[["gini_index","unemployment"]]
model2x =
x_train[["mortality_rate_infant","intentional_homicides","suicide_mortality_rate"]]
model3x =
x_train[["government_health_expenditure_pct_gdp","expenditure_on_education_pct_gdp"]]
model4x = x_train[["population_density","urban_population_pct_total"]]
model5x =
x_train[["tax_revenue","taxes_on_income_profits_capital","gdp_per_capi
```

```

ta_ppp"]]]
model6x =
x_train[["gini_index","mortality_rate_infant","intentional_homicides",
"suicide_mortality_rate", \

"alcohol_consumption_per_capita","government_health_expenditure_pct_gd
p"]]]
model7x =
x_train[["alcohol_consumption_per_capita","intentional_homicides","urb
an_population_pct_total"]]

# calcualtes best RIDGE and LASSO alphas for each model using cross-
validation from training dataset
# RIDGE
ridgeCV = RidgeCV(alphas =
np.arange(0.01,100,0.05)).fit(model1x,y_train)
alpha1r = ridgeCV.alpha_
print("Most appropriate Ridge alpha for model 1:", alpha1r)

ridgeCV = RidgeCV(alphas =
np.arange(0.01,100,0.05)).fit(model2x,y_train)
alpha2r = ridgeCV.alpha_
print("Most appropriate Ridge alpha for model 2:", alpha2r)

ridgeCV = RidgeCV(alphas =
np.arange(0.01,100,0.05)).fit(model3x,y_train)
alpha3r = ridgeCV.alpha_
print("Most appropriate Ridge alpha for model 3:", alpha3r)

ridgeCV = RidgeCV(alphas =
np.arange(0.01,100,0.05)).fit(model4x,y_train)
alpha4r = ridgeCV.alpha_
print("Most appropriate Ridge alpha for model 4:", alpha4r)

ridgeCV = RidgeCV(alphas =
np.arange(0.01,100,0.05)).fit(model5x,y_train)
alpha5r = ridgeCV.alpha_
print("Most appropriate Ridge alpha for model 5:", alpha5r)

ridgeCV = RidgeCV(alphas =
np.arange(0.01,100,0.05)).fit(model6x,y_train)
alpha6r = ridgeCV.alpha_
print("Most appropriate Ridge alpha for model 6:", alpha6r)

ridgeCV = RidgeCV(alphas =
np.arange(0.01,100,0.05)).fit(model7x,y_train)
alpha7r = ridgeCV.alpha_
print("Most appropriate Ridge alpha for model 7:", alpha7r, "\n")

```

```

# LASSO
lassoCV = LassoCV(cv = None, n_alphas = 100).fit(model1x,y_train)
alpha1l = lassoCV.alpha_
print("Most appropriate LASSO alpha for model 1:", alpha1l)

lassoCV = LassoCV(cv = None, n_alphas = 100).fit(model2x,y_train)
alpha2l = lassoCV.alpha_
print("Most appropriate LASSO alpha for model 2:", alpha2l)

lassoCV = LassoCV(cv = None, n_alphas = 100).fit(model3x,y_train)
alpha3l = lassoCV.alpha_
print("Most appropriate LASSO alpha for model 3:", alpha3l)

lassoCV = LassoCV(cv = None, n_alphas = 100).fit(model4x,y_train)
alpha4l = lassoCV.alpha_
print("Most appropriate LASSO alpha for model 4:", alpha4l)

lassoCV = LassoCV(cv = None, n_alphas = 100).fit(model5x,y_train)
alpha5l = lassoCV.alpha_
print("Most appropriate LASSO alpha for model 5:", alpha5l)

lassoCV = LassoCV(cv = None, n_alphas = 100).fit(model6x,y_train)
alpha6l = lassoCV.alpha_
print("Most appropriate LASSO alpha for model 6:", alpha6l)

lassoCV = LassoCV(cv = None, n_alphas = 100).fit(model7x,y_train)
alpha7l = lassoCV.alpha_
print("Most appropriate LASSO alpha for model 7:", alpha7l)

Most appropriate Ridge alpha for model 1: 5.96
Most appropriate Ridge alpha for model 2: 9.01
Most appropriate Ridge alpha for model 3: 99.960000000000001
Most appropriate Ridge alpha for model 4: 92.710000000000001
Most appropriate Ridge alpha for model 5: 79.26
Most appropriate Ridge alpha for model 6: 25.160000000000004
Most appropriate Ridge alpha for model 7: 1.01

Most appropriate LASSO alpha for model 1: 0.009035105986016721
Most appropriate LASSO alpha for model 2: 0.6709563385389306
Most appropriate LASSO alpha for model 3: 2.473217768450261
Most appropriate LASSO alpha for model 4: 0.033625551095921184
Most appropriate LASSO alpha for model 5: 1.2976570742114053
Most appropriate LASSO alpha for model 6: 0.4414444674193534
Most appropriate LASSO alpha for model 7: 0.007559802997655883

# implementing K-fold to test for best model (manual)

# initializes k group IDs, creating modified training dataframes for
cross-validation
kgroup_ids = np.arange(1,6)

```

```

kgroups_array = np.repeat(kgroup_ids,27)
np.random.shuffle(kgroups_array)
x_train["k_group"] = kgroups_array
y_train_cv= pd.DataFrame({'rating':y_train,'k_group':kgroups_array})

y_train1 = y_train_cv.loc[y_train_cv["k_group"] != 1]['rating']

# implementing Kfold to select the best model
ridge_1_mses = []
ridge_2_mses = []
ridge_3_mses = []
ridge_4_mses = []
ridge_5_mses = []
ridge_6_mses = []
ridge_7_mses = []

linear_1_mses = []
linear_2_mses = []
linear_3_mses = []
linear_4_mses = []
linear_5_mses = []
linear_6_mses = []
linear_7_mses = []

lasso_1_mses = []
lasso_2_mses = []
lasso_3_mses = []
lasso_4_mses = []
lasso_5_mses = []
lasso_6_mses = []
lasso_7_mses = []

best_r2s = []

for k in kgroup_ids:
    print('working on fold', k, "...")
    x_train1 = x_train.loc[x_train["k_group"] != k]
    y_train1 = y_train_cv.loc[y_train_cv["k_group"] != k]['rating']

    model1x = x_train1[["gini_index","unemployment"]]
    model2x =
x_train1[["mortality_rate_infant","intentional_homicides","suicide_mor
tality_rate"]]
    model3x =
x_train1[["government_health_expenditure_pct_gdp","expenditure_on
_education_pct_gdp"]]
    model4x =
x_train1[["population_density","urban_population_pct_total"]]
    model5x =

```

```

x_train1[["tax_revenue", "taxes_on_income_profits_capital", "gdp_per_capita_ppp"]]
model6x =
x_train1[["gini_index", "mortality_rate_infant", "intentional_homicides",
"suicide_mortality_rate", \

"alcohol_consumption_per_capita", "government_health_expenditure_pct_gdp"]]
model7x =
x_train1[["alcohol_consumption_per_capita", "intentional_homicides", "urban_population_pct_total"]]

model1tst = x_test[["gini_index", "unemployment"]]
model2tst =
x_test[["mortality_rate_infant", "intentional_homicides", "suicide_mortality_rate"]]
model3tst =
x_test[["government_health_expenditure_pct_gdp", "expenditure_on_education_pct_gdp"]]
model4tst =
x_test[["population_density", "urban_population_pct_total"]]
model5tst =
x_test[["tax_revenue", "taxes_on_income_profits_capital", "gdp_per_capita_ppp"]]
model6tst =
x_test[["gini_index", "mortality_rate_infant", "intentional_homicides", "suicide_mortality_rate", \

"alcohol_consumption_per_capita", "government_health_expenditure_pct_gdp"]]
model7tst =
x_test[["alcohol_consumption_per_capita", "intentional_homicides", "urban_population_pct_total"]]

ridgereg = Ridge(alpha = alpha1r).fit(model1x, y_train1)
predicted_y1 = ridgereg.predict(model1tst)
print("-----")
# print("MSE for Model 1 Ridge regression:", mean_squared_error(y_test, predicted_y1))
ridge_1_mses.append(mean_squared_error(y_test, predicted_y1))

ridgereg = Ridge(alpha = alpha2r).fit(model2x, y_train1)
predicted_y2 = ridgereg.predict(model2tst)
# print("MSE for Model 2 Ridge regression:", mean_squared_error(y_test, predicted_y2))
ridge_2_mses.append(mean_squared_error(y_test, predicted_y2))

ridgereg = Ridge(alpha = alpha3r).fit(model3x, y_train1)
predicted_y3 = ridgereg.predict(model3tst)

```

```

# print("MSE for Model 3 Ridge
regression:",mean_squared_error(y_test,predicted_y3))
ridge_3_mses.append(mean_squared_error(y_test,predicted_y3))

ridgereg = Ridge(alpha = alpha4r).fit(model4x,y_train1)
predicted_y4 = ridgereg.predict(model4tst)
# print("MSE for Model 4 Ridge
regression:",mean_squared_error(y_test,predicted_y4))
ridge_4_mses.append(mean_squared_error(y_test,predicted_y4))

ridgereg = Ridge(alpha = alpha5r).fit(model5x,y_train1)
predicted_y5 = ridgereg.predict(model5tst)
# print("MSE for Model 5 Ridge
regression:",mean_squared_error(y_test,predicted_y5))
ridge_5_mses.append(mean_squared_error(y_test,predicted_y5))

ridgereg = Ridge(alpha = alpha6r).fit(model6x,y_train1)
predicted_y6 = ridgereg.predict(model6tst)
# print("MSE for Model 6 Ridge
regression:",mean_squared_error(y_test,predicted_y6))
ridge_6_mses.append(mean_squared_error(y_test,predicted_y6))
best_r2s.append(r2_score(y_test,predicted_y6))

ridgereg = Ridge(alpha = alpha7r).fit(model7x,y_train1)
predicted_y7 = ridgereg.predict(model7tst)
# print("MSE for Model 7 Ridge
regression:",mean_squared_error(y_test,predicted_y7))
ridge_7_mses.append(mean_squared_error(y_test,predicted_y7))

linearReg = LinearRegression().fit(model1x,y_train1)
predicted_y1 = linearReg.predict(model1tst)
# print("MSE for Model 1 Linear
regression:",mean_squared_error(y_test,predicted_y1))
linear_1_mses.append(mean_squared_error(y_test,predicted_y1))

linearReg = LinearRegression().fit(model2x,y_train1)
predicted_y2 = linearReg.predict(model2tst)
# print("MSE for Model 2 Linear
regression:",mean_squared_error(y_test,predicted_y2))
linear_2_mses.append(mean_squared_error(y_test,predicted_y2))

linearReg = LinearRegression().fit(model3x,y_train1)
predicted_y3 = linearReg.predict(model3tst)
# print("MSE for Model 3 Linear
regression:",mean_squared_error(y_test,predicted_y3))
linear_3_mses.append(mean_squared_error(y_test,predicted_y3))

linearReg = LinearRegression().fit(model4x,y_train1)
predicted_y4 = linearReg.predict(model4tst)
# print("MSE for Model 4 Linear

```

```

regression:",mean_squared_error(y_test,predicted_y4))
linear_4_mses.append(mean_squared_error(y_test,predicted_y4))

linearReg = LinearRegression().fit(model5x,y_train1)
predicted_y5 = linearReg.predict(model5tst)
# print("MSE for Model 5 Linear
regression:",mean_squared_error(y_test,predicted_y5))
linear_5_mses.append(mean_squared_error(y_test,predicted_y5))

linearReg = LinearRegression().fit(model6x,y_train1)
predicted_y6 = linearReg.predict(model6tst)
# print("MSE for Model 6 Linear
regression:",mean_squared_error(y_test,predicted_y6))
linear_6_mses.append(mean_squared_error(y_test,predicted_y6))

linearReg = LinearRegression().fit(model7x,y_train1)
predicted_y7 = linearReg.predict(model7tst)
# print("MSE for Model 7 Linear
regression:",mean_squared_error(y_test,predicted_y7))
linear_7_mses.append(mean_squared_error(y_test,predicted_y7))

lassoreg = Lasso(alpha = alpha1l).fit(model1x,y_train1)
predicted_y1 = lassoreg.predict(model1tst)
# print("MSE for Model 1 LASSO
regression:",mean_squared_error(y_test,predicted_y1))
lasso_1_mses.append(mean_squared_error(y_test,predicted_y1))

lassoreg = Lasso(alpha = alpha2l).fit(model2x,y_train1)
predicted_y2 = lassoreg.predict(model2tst)
# print("MSE for Model 2 LASSO
regression:",mean_squared_error(y_test,predicted_y2))
lasso_2_mses.append(mean_squared_error(y_test,predicted_y2))

lassoreg = Lasso(alpha = alpha3l).fit(model3x,y_train1)
predicted_y3 = lassoreg.predict(model3tst)
# print("MSE for Model 3 LASSO
regression:",mean_squared_error(y_test,predicted_y3))
lasso_3_mses.append(mean_squared_error(y_test,predicted_y3))

lassoreg = Lasso(alpha = alpha4l).fit(model4x,y_train1)
predicted_y4 = lassoreg.predict(model4tst)
# print("MSE for Model 4 LASSO
regression:",mean_squared_error(y_test,predicted_y4))
lasso_4_mses.append(mean_squared_error(y_test,predicted_y4))

lassoreg = Lasso(alpha = alpha5l).fit(model5x,y_train1)
predicted_y5 = lassoreg.predict(model5tst)
# print("MSE for Model 5 LASSO
regression:",mean_squared_error(y_test,predicted_y5))
lasso_5_mses.append(mean_squared_error(y_test,predicted_y5))

```



```

lassoreg = Lasso(alpha = alpha6l).fit(model6x,y_train1)
predicted_y6 = lassoreg.predict(model6tst)
# print("MSE for Model 6 LASSO
regression:",mean_squared_error(y_test,predicted_y6))
lasso_6_mses.append(mean_squared_error(y_test,predicted_y6))

lassoreg = Lasso(alpha = alpha7l).fit(model7x,y_train1)
predicted_y7 = lassoreg.predict(model7tst)
# print("MSE for Model 7 LASSO
regression:",mean_squared_error(y_test,predicted_y7))
lasso_7_mses.append(mean_squared_error(y_test,predicted_y7))
print(' ')

print("Ridge Model 1 Mean MSE:",np.mean(ridge_1_mses))
print("Ridge Model 2 Mean MSE:",np.mean(ridge_2_mses))
print("Ridge Model 3 Mean MSE:",np.mean(ridge_3_mses))
print("Ridge Model 4 Mean MSE:",np.mean(ridge_4_mses))
print("Ridge Model 5 Mean MSE:",np.mean(ridge_5_mses))
print("Ridge Model 6 Mean MSE:",np.mean(ridge_6_mses))
print("Ridge Model 7 Mean MSE:",np.mean(ridge_7_mses))
print(" ")
print("Linear Model 1 Mean MSE:",np.mean(linear_1_mses))
print("Linear Model 2 Mean MSE:",np.mean(linear_2_mses))
print("Linear Model 3 Mean MSE:",np.mean(linear_3_mses))
print("Linear Model 4 Mean MSE:",np.mean(linear_4_mses))
print("Linear Model 5 Mean MSE:",np.mean(linear_5_mses))
print("Linear Model 6 Mean MSE:",np.mean(linear_6_mses))
print("Linear Model 7 Mean MSE:",np.mean(linear_7_mses))
print(" ")
print("LASSO Model 1 Mean MSE:",np.mean(lasso_1_mses))
print("LASSO Model 2 Mean MSE:",np.mean(lasso_2_mses))
print("LASSO Model 3 Mean MSE:",np.mean(lasso_3_mses))
print("LASSO Model 4 Mean MSE:",np.mean(lasso_4_mses))
print("LASSO Model 5 Mean MSE:",np.mean(lasso_5_mses))
print("LASSO Model 6 Mean MSE:",np.mean(lasso_6_mses))
print("LASSO Model 7 Mean MSE:",np.mean(lasso_7_mses))
print(" ")

print("Best Model (Ridge 6) Mean R^2:",np.mean(best_r2s))

working on fold 1 ...
-----

working on fold 2 ...
-----

working on fold 3 ...
-----

```

```
working on fold 4 ...
```

```
-----
```

```
working on fold 5 ...
```

```
-----
```

```
Ridge Model 1 Mean MSE: 348.030073368065  
Ridge Model 2 Mean MSE: 345.24193578765033  
Ridge Model 3 Mean MSE: 431.8989196823908  
Ridge Model 4 Mean MSE: 429.2027447216643  
Ridge Model 5 Mean MSE: 413.136386745325  
Ridge Model 6 Mean MSE: 336.05341183074245  
Ridge Model 7 Mean MSE: 409.1743014925128
```

```
Linear Model 1 Mean MSE: 349.6230579375316  
Linear Model 2 Mean MSE: 345.6865492881978  
Linear Model 3 Mean MSE: 426.97797907901924  
Linear Model 4 Mean MSE: 430.2551745905697  
Linear Model 5 Mean MSE: 419.63504795191665  
Linear Model 6 Mean MSE: 337.429379618244  
Linear Model 7 Mean MSE: 410.68525388631633
```

```
LASSO Model 1 Mean MSE: 349.5833427854218  
LASSO Model 2 Mean MSE: 348.9848770294435  
LASSO Model 3 Mean MSE: 435.8617145144079  
LASSO Model 4 Mean MSE: 430.2436627504747  
LASSO Model 5 Mean MSE: 407.10899545782615  
LASSO Model 6 Mean MSE: 336.32302023314764  
LASSO Model 7 Mean MSE: 410.55284349573355
```

```
Best Model (Ridge 6) Mean R^2: 0.22449780296456773
```

```
# recursively implementing K-fold to test for best model a number of  
times  
# will contextualize how the ORIGINAL test-train split affects best  
model
```

```
# a list of the 7 chosen possible models, in order model 1-7
```

```
models = [["gini_index", "unemployment"],
```

```
["mortality_rate_infant", "intentional_homicides", "suicide_mortality_ra  
te"],
```

```
["government_health_expenditure_pct_gdp", "expenditure_on  
_education_pct_gdp"],
```

```
["population_density", "urban_population_pct_total"],
```

```
["tax_revenue", "taxes_on_income_profits_capital", "gdp_per_capita_ppp"]
```

```
,
```

```
["gini_index", "mortality_rate_infant", "intentional_homicides", "suicide
```

```

_mortality_rate", \
"alcohol_consumption_per_capita","government_health_expenditure_pct_gdp"],

["alcohol_consumption_per_capita","intentional_homicides","urban_population_pct_total"]]

# for each K-fold CV, records best model (out of 21), and their MSEs and R2s
best_models = np.array([])
best_mses = np.array([])
best_r2s = np.array([])

# THIS PART WILL TAKE LONG TO RUN: CAN LIMIT K-FOLDS TO ONE:
range(100) -> range(1)
for split in notebook.tqdm(range(100)):
    # creates a shuffled K-fold indices of given folds (default=5)
    k = 5
    kfold = KFold(n_splits=k,shuffle=True)

    model_mses = np.zeros((3,7))
    model_r2s = np.zeros((3,7))

    # does a randomized initial test-train split before running K-Fold CV
    x_traink, x_testk, y_traink, y_testk = train_test_split(x_df, y_df,
test_size = 0.2)

    # iterates over each of the models
    for m,model in enumerate(models):
        # calculates best fitting RIDGE and LASSO alphas using new initial
training dataframe and model
        ridgeCV = RidgeCV(alphas =
np.arange(0.01,100,0.05)).fit(x_traink[model],y_traink)
        ralpha = ridgeCV.alpha_
        lassoCV =
LassoCV(cv=None,n_alphas=500).fit(x_traink[model],y_traink)
        lalpha = lassoCV.alpha_

        # for each K-fold, record that fold's MSE and R2 values
        mses = np.zeros((3,k))
        r2s = np.zeros((3,k))
        # fold number (0 to k-1)
        ki = 0

        # performs K-fold CV
        for train_index,test_index in kfold.split(x_traink):
            # test and train split for specific fold, from initial training
dataframe and model

```

```

X_tr, X_te, y_tr, y_te = x_traink.iloc[train_index],
x_traink.iloc[test_index], y_traink.iloc[train_index],
y_traink.iloc[test_index]

# performs linear regression on training fold, predicts y using
test fold
linearReg = LinearRegression().fit(X_tr[model],y_tr)
predicted_y = linearReg.predict(X_te[model])
mses[0][ki] = mean_squared_error(y_te,predicted_y)
r2s[0][ki] = r2_score(y_te,predicted_y)

# performs RIDGE regression on training fold, predicts y using
test fold
ridgeReg = Ridge(alpha = ralpha).fit(X_tr[model],y_tr)
predicted_y = ridgeReg.predict(X_te[model])
mses[1][ki] = mean_squared_error(y_te,predicted_y)
r2s[1][ki] = r2_score(y_te,predicted_y)

# performs LASSO regression on training fold, predicts y using
test fold
lassoReg = Lasso(alpha = lalpha).fit(X_tr[model],y_tr)
predicted_y = lassoReg.predict(X_te[model])
mses[2][ki] = mean_squared_error(y_te,predicted_y)
r2s[2][ki] = r2_score(y_te,predicted_y)

# increases fold number
ki += 1

# across all folds, calculates the mean MSE
model_mses[0][m] = np.mean(mses[0])
model_mses[1][m] = np.mean(mses[1])
model_mses[2][m] = np.mean(mses[2])

# across all folds, calculates the mean R2 score
model_r2s[0][m] = np.mean(r2s[0])
model_r2s[1][m] = np.mean(r2s[1])
model_r2s[2][m] = np.mean(r2s[2])

# out of a given K-Fold CV, appends best model, and its MSE and R2
best_models = np.append(best_models,np.argmin(model_mses))
best_mses = np.append(best_mses,np.min(model_mses))
best_r2s = np.append(best_r2s,model_r2s[np.argmin(model_mses)//7]
[np.argmin(model_mses)%7])

print("Best Models (#):",best_models)
print("Best Model MSE:",best_mses)
print("Best Model R^2:",best_r2s)

{"model_id":"f67600c6e2a24411b44f762096f7a4b4","version_major":2,"version_minor":0}

```

Best Models (#): [7. 8. 12. 5. 1. 7. 12. 12. 12. 12. 12. 0. 7.
12. 5. 19. 12. 19.
0. 19. 7. 7. 12. 12. 1. 19. 7. 7. 12. 12. 8. 12. 12. 12. 12.
12.
12. 19. 7. 12. 7. 12. 0. 12. 12. 12. 19. 12. 7. 12. 0. 5. 5.
7.
12. 12. 19. 12. 12. 12. 12. 8. 7. 7. 7. 12. 12. 1. 12. 19. 12.
7.
7. 5. 12. 12. 19. 12. 19. 1. 12. 12. 12. 12. 7. 12. 12. 8. 12.
12.
12. 12. 8. 1. 7. 12. 7. 7. 0. 0.]

Best Model MSE: [324.26169157 301.59536948 334.54295436 333.86602521
310.92577076

306.30604768 316.23188624 293.21461502 340.50531055 332.49428351
335.95854873 318.88896443 308.93343671 328.70345973 323.51321629
304.90838875 309.40546515 327.60486263 329.47162582 312.27507599
350.08862812 312.5815229 344.80496493 335.17902564 318.08228269
313.45059923 327.31112444 327.4336232 359.59052334 332.26992543
313.75529574 290.27651095 303.40757796 270.57183917 337.1653541
330.39562629 347.3398795 285.00974694 311.95726817 338.31546224
321.16356575 331.27535585 296.45381751 329.64092784 289.74159367
339.69579174 334.82136214 291.66169041 293.58286429 302.70480685
320.06964839 313.40873098 294.32080421 345.54541267 324.69376476
299.48611478 337.68235155 297.06806331 337.19769647 316.38276524
281.8228077 340.41877254 320.13909966 287.36115712 320.8692162
319.99187508 320.35495003 334.2169424 330.30596196 300.02718358
327.53752575 317.60599811 309.37904528 306.4347881 324.34741628
322.79389878 322.39098658 314.11439445 323.3793774 319.53797618
329.92853867 336.02640699 320.51645237 325.09582641 307.49531227
308.03975115 337.36720691 327.92953415 314.89601368 312.82864106
324.49649328 314.07103748 307.6356152 335.70545744 334.17883775
320.62120772 327.8438504 304.80396453 296.85030683 303.2063172]

Best Model R^2: [0.23524283 0.2329234 0.26902664 0.12564711
0.25838466 0.2501753

0.25145202 0.29762884 0.21266328 0.18612768 0.28442602 0.21976035
0.21567122 0.26907972 0.15098136 0.25088636 0.14663318 0.22595114
0.25708259 0.18087079 0.25731593 0.17525782 0.14865546 0.25698496
0.21448967 0.28392619 0.14821688 0.17711886 0.23484425 0.26908705
0.24128525 0.29398777 0.32641148 0.30762301 0.22297186 0.17498078
0.18712769 0.27439243 0.19650276 0.26639014 0.15349274 0.28822626
0.23080729 0.19418514 0.23182475 0.21905316 0.28523682 0.26388005
0.22457687 0.23679096 0.17680572 0.18704896 0.29740885 0.1210857
0.21281665 0.25656137 0.27175736 0.16636146 0.11817928 0.16521081
0.29335891 0.19197067 0.19289006 0.19505283 0.18714827 0.25567163
0.24258664 0.15529273 0.21021812 0.35429824 0.26605735 0.21793912
0.17351692 0.21634219 0.18678703 0.19649822 0.24395725 0.23622035
0.12282165 0.22658188 0.22386951 0.26508616 0.28652673 0.16079967
0.1543984 0.20087495 0.184429 0.165907 0.11644778 0.27728154

```
0.23937888 0.27133759 0.18687607 0.21579371 0.19338743 0.30977655
0.23577436 0.22451223 0.30783951 0.22631346]
```

```
# results from recursive K-Fold tries
```

```
# matches model number to an informative name
```

```
models_s = []
```

```
for mo in range(21):
```

```
    if mo//7==0:
```

```
        models_s.append("model"+str(mo%7+1)+"_linear")
```

```
    elif mo//7==1:
```

```
        models_s.append("model"+str(mo%7+1)+"_ridge")
```

```
    else:
```

```
        models_s.append("model"+str(mo%7+1)+"_lasso")
```

```
# gets best models, and their occurances over the number of K-Folds
```

```
x,height = np.unique(best_models,return_counts=True)
```

```
best_models_df = pd.DataFrame({"model":x,"count":height})
```

```
best_models_df["model"] =
```

```
best_models_df["model"].replace(np.arange(21),models_s)
```

```
# plots histogram of best model over multiple K-Folds
```

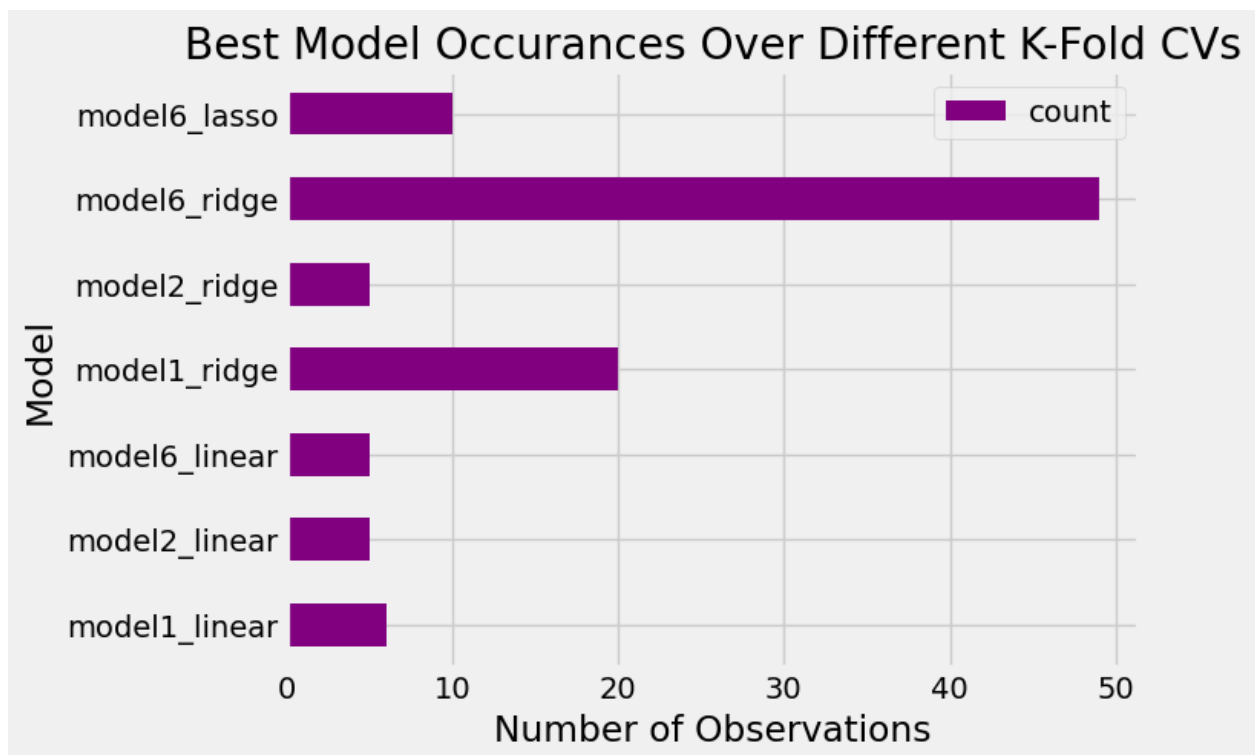
```
best_models_df.plot.barh(x="model",y="count",color="purple")
```

```
plt.title("Best Model Occurances Over Different K-Fold CVs")
```

```
plt.ylabel("Model")
```

```
plt.xlabel("Number of Observations")
```

```
plt.show()
```



```

# choosing best model based on histogram; running a normal CV using
original test-train split

# best model: model 6, RIDGE
model6_train =
x_train[["gini_index","mortality_rate_infant","intentional_homicides",
"suicide_mortality_rate", \

"alcohol_consumption_per_capita","government_health_expenditure_pct_gd
p"]]
model6_test =
x_test[["gini_index","mortality_rate_infant","intentional_homicides",
"suicide_mortality_rate", \

"alcohol_consumption_per_capita","government_health_expenditure_pct_gd
p"]]

# recalculates and saves RIDGE alpha for clarity
best_ridgeCV = RidgeCV(alphas =
np.arange(0.01,100,0.05)).fit(model6_train,y_train)
best_alpha = best_ridgeCV.alpha_

# performs regression, predicts using test dataframe
best_ridge = Ridge(alpha=best_alpha).fit(model6_train,y_train)
best_predicted_y = best_ridge.predict(model6_test)

# calculates MSE and r2
best_mse = mean_squared_error(y_test,predicted_y6)
best_r2 = r2_score(y_test,predicted_y6)

```

BEST MODEL ANALYSIS

```

# looking at the best model's residuals & comparing when we also
control for country
print("MSE for Model 6 Ridge regression:",best_mse)

# plotting the residuals for the best model
plt.scatter(best_predicted_y,best_predicted_y - y_test)
plt.ylim(-41,41)
plt.axhline(0,color='red',linewidth = 2)
plt.title('Residuals Plot: Best Model')
plt.xlabel("Predicted Values")
plt.ylabel("Residuals")
plt.show()

# getting the MSE for the best model + controlling for country
country_control =
x_train[["gini_index","mortality_rate_infant","intentional_homicides",
"suicide_mortality_rate", \

```

```

"alcohol_consumption_per_capita","government_health_expenditure_pct_gd
p",'country_AUT', 'country_CAN', \
    'country_CHE', 'country_CZE','country_DEU',
'country_DNK', 'country_ESP', 'country_EST','country_FIN', \
    'country_FRA', 'country_GBR',
'country_GRC','country_HUN', 'country_IRL', 'country_ISL',
'country_ISR','country_ITA', \
    'country_KOR', 'country_LUX',
'country_LVA','country_NLD', 'country_NOR', 'country_POL',
'country_PRT','country_SVK', \
    'country_SVN', 'country_SWE',
'country_USA']]

```

```

ridgeCV =
RidgeCV(alphas=np.arange(0.01,100,0.05)).fit(country_control,y_train)
print("Most appropriate Ridge alpha for model 6 (with country
control):", ridgeCV.alpha_)

```

```

country_control_tst =
x_test[["gini_index","mortality_rate_infant","intentional_homicides","
suicide_mortality_rate", \

```

```

"alcohol_consumption_per_capita","government_health_expenditure_pct_gd
p",'country_AUT', 'country_CAN', \
    'country_CHE', 'country_CZE','country_DEU',
'country_DNK', 'country_ESP', 'country_EST','country_FIN', \
    'country_FRA', 'country_GBR',
'country_GRC','country_HUN', 'country_IRL', 'country_ISL',
'country_ISR','country_ITA', \
    'country_KOR', 'country_LUX',
'country_LVA','country_NLD', 'country_NOR', 'country_POL',
'country_PRT','country_SVK', \
    'country_SVN', 'country_SWE',
'country_USA']]

```

```

ridgereg = Ridge(alpha=ridgeCV.alpha_).fit(country_control,y_train)
predicted_y_cc = ridgereg.predict(country_control_tst)
print("MSE for Model 6 + Country Variable Ridge
regression:",mean_squared_error(y_test,predicted_y_cc))


```

```

plt.scatter(predicted_y_cc,predicted_y_cc-y_test)
plt.axhline(0,color='red',linewidth = 2)
plt.ylim(-41,41)
plt.title('Residuals Plot: Best Model + Country Variable')
plt.xlabel("Predicted Values")
plt.ylabel("Residuals")
plt.show()


```

MSE for Model 6 Ridge regression: 337.1024391863911



Most appropriate Ridge alpha for model 6 (with country control): 0.26
MSE for Model 6 + Country Variable Ridge regression:
106.16160153621229

Residuals Plot: Best Model + Country Variable




```
# analyze how countries cluster over influential variables and the
PISA score
# example: suicide_mortality_rate


colors =
["b","g","r","c","m","y","k","hotpink","orangered","gold","teal","viol
et","maroon","slategrey","blueviolet","peru","yellowgreen" \
    ,"navy","slateblue","rebeccapurple","lightblue","seagreen",
rosybrown","dodgerblue","khaki","darkgreen","lightcoral","darkslategre
y","limegreen"]

for i,country in enumerate(pisa_df_tot.country.unique()):
    plt.scatter(pisa_df_tot.loc[pisa_df_tot["country"] ==
country].suicide_mortality_rate,pisa_df_tot.loc[pisa_df_tot["country"]
== country].rating, \
                label = country,color = colors[i])
    plt.legend(ncols=5,loc='upper center',bbox_to_anchor=(0.5, -0.11))
    plt.xlabel("(standardized) suicide_mortality_rate")
    plt.ylabel("PISA Rating")
plt.show()

print("\n")
for i,country in enumerate(pisa_df_tot.country.unique()):
```

```
plt.scatter(pisa_df_tot.loc[pisa_df_tot["country"] ==
country].unemployment,pisa_df_tot.loc[pisa_df_tot["country"] ==
country].rating, \
            label = country,color = colors[i])
plt.legend(ncols=5,loc='upper center',bbox_to_anchor=(0.5, -0.11))
plt.xlabel("(standardized) unemployment")
plt.ylabel("PISA Rating")
plt.show()
```





CLUSTERING ANALYSIS

```
# define dataframe for clustering, with best model including the PISA
score
bestmodel_df =
pisa_df_tot[["gini_index", "mortality_rate_infant", "intentional_homicid
es", \

"suicide_mortality_rate", "alcohol_consumption_per_capita", "government_
health_expenditure_pct_gdp", "rating"]]

# standardize PISA score
scaler = StandardScaler()
bestmodel_df[["rating"]] =
scaler.fit_transform(bestmodel_df[["rating"]])
```

```
# display the dataframe
```

```
display(bestmodel_df.head())
```

```
C:\Users\daniel\AppData\Local\Temp\ipykernel_22508\2792784023.py:7:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation:
```

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#  
returning-a-view-versus-a-copy
```

```
bestmodel_df[["rating"]] =  
scaler.fit_transform(bestmodel_df[["rating"]])
```

	gini_index	mortality_rate_infant	
intentional_homicides \			
index_code			
AUS-2003	0.462126	0.831882	0.142289
AUS-2006	0.559458	0.679648	-0.005489
AUS-2009	0.559458	0.299063	-0.152009
AUS-2012	0.559458	-0.157639	-0.285880
AUS-2015	0.559458	-0.385990	-0.358186

	suicide_mortality_rate	alcohol_consumption_per_capita	\
index_code			
AUS-2003	-0.660127	0.121143	
AUS-2006	-0.642575	0.121143	
AUS-2009	-0.537262	0.121143	
AUS-2012	-0.449501	0.121143	
AUS-2015	-0.186219	0.121143	

	government_health_expenditure_pct_gdp	rating
index_code		
AUS-2003	-0.330309	1.206663
AUS-2006	-0.269865	1.015228
AUS-2009	0.059373	0.728076
AUS-2012	0.018771	0.249488
AUS-2015	0.681248	-0.229099

```
# required definitions to perform K-Means Clustering
```

```
def distance(pt1, pt2):
```

```
    """Return the distance between two points, represented as  
    arrays"""
```

```

    return np.sqrt(sum((pt1 - pt2)**2))

def initialize_centroids(df,K):
    random_ids = np.random.permutation(df.shape[0])
    centroids = df.iloc[random_ids[:K]]
    return centroids

def compute_distance(df, centroids):
    K=centroids.shape[0]
    distances_ar = np.zeros((df.shape[0], K))
    for k in range(K):
        point=centroids.iloc[k]
        def distance_from_point(row):
            return distance(point, np.array(row))
        distances_ar[:,k] =
df.apply(distance_from_point,axis=1).values
    return distances_ar

def compute_sse(df, labels, centroids,K):
    distances_ar = np.zeros(df.shape[0])
    for k in range(K):
        point=centroids.iloc[k]
        def distance_from_point(row):
            return distance(point, np.array(row))
        distances_ar[labels == k] = df[labels ==
k].apply(distance_from_point,axis=1).values
    return np.sum(distances_ar)

def compute_centroids(df, labels, K):
    centroids = np.zeros((K, df.shape[1]))
    for k in range(K):
        centroids[k, :] = df[labels == k].mean()
    return centroids

def Kmeans(df,K):
    max_iter=20

    centroids=initialize_centroids(df,K)

    for i in range(max_iter):
        old_centroids = centroids
        dist_matrix = compute_distance(df, old_centroids)
        clust=np.argmin(dist_matrix,axis = 1)
        centroids = pd.DataFrame(compute_centroids(df,clust,K))

    return centroids,clust

def Kmeans_sse(df,K):
    '''performs Kmeans returns centroids and prints sse of each new
centroids'''

```

```

#define the maximum number of iterations
max_iter=20

#initialize centroids
centroids=initialize_centroids(df,K)

for i in range(max_iter):
    old_centroids = centroids
    dist_matrix = compute_distance(df, old_centroids)
    clust=np.argmin(dist_matrix, axis=1)
    centroids = pd.DataFrame(compute_centroids(df,clust,K))

# return the centroids
return compute_sse(df,clust,old_centroids,K)

# create k and SSE arrays to find best number of clusters

k = np.array([])
sse = np.array([])
for i in np.arange(1,10):
    k = np.append(k,i)
    sse = np.append(sse,Kmeans_sse(bestmodel_df,i))

# elbow plot is used to select number of clusters to use

plt.plot(k,sse)
plt.title("Elbow Plot")
plt.xlabel("K (number of clusters)")
plt.ylabel("Minimum SSE")
plt.show()

```



```
# return centroids and clusters used chosen K=3
centroids, clust = Kmeans(bestmodel_df, 3)

# show the clusters created by K-Means Clustering!
pd.set_option('display.max_rows', None)

country_names = bestmodel_df.index
cluster_countries =
pd.DataFrame({"country": country_names, "cluster_label": clust})
cluster_countries

cluster0_countries =
cluster_countries[cluster_countries["cluster_label"] == 0]
display(cluster0_countries)

cluster1_countries =
cluster_countries[cluster_countries["cluster_label"] == 1]
display(cluster1_countries)

cluster2_countries =
```



```
cluster_countries[cluster_countries["cluster_label"] == 2]  
display(cluster2_countries)
```

	country	cluster_label
6	AUT-2003	0
17	CZE-2003	0
18	CZE-2006	0
19	CZE-2009	0
20	CZE-2012	0
21	CZE-2015	0
22	CZE-2018	0
29	EST-2006	0
30	EST-2009	0
31	EST-2012	0
32	EST-2015	0
33	EST-2018	0
34	FIN-2003	0
35	FIN-2006	0
36	FIN-2009	0
58	HUN-2003	0
59	HUN-2006	0
60	HUN-2009	0
61	HUN-2012	0
62	HUN-2015	0
63	HUN-2018	0
87	KOR-2003	0
88	KOR-2006	0
89	KOR-2009	0
90	KOR-2012	0
91	KOR-2015	0
92	KOR-2018	0
93	LVA-2003	0
94	LVA-2006	0
95	LVA-2009	0
96	LVA-2012	0
97	LVA-2015	0
98	LVA-2018	0
99	LUX-2003	0
100	LUX-2006	0
117	POL-2003	0
118	POL-2006	0
119	POL-2009	0
120	POL-2012	0
121	POL-2015	0
122	POL-2018	0
129	SVK-2003	0
130	SVK-2006	0
131	SVK-2009	0
132	SVK-2012	0
134	SVK-2018	0

135	SVN-2006	0
136	SVN-2009	0
137	SVN-2012	0
138	SVN-2015	0
139	SVN-2018	0
152	CHE-2003	0
153	CHE-2006	0
154	CHE-2009	0
155	CHE-2012	0
156	CHE-2015	0
157	CHE-2018	0

	country	cluster_label
0	AUS-2003	1
1	AUS-2006	1
2	AUS-2009	1
3	AUS-2012	1
4	AUS-2015	1
5	AUS-2018	1
7	AUT-2006	1
8	AUT-2012	1
9	AUT-2015	1
10	AUT-2018	1
11	CAN-2003	1
12	CAN-2006	1
13	CAN-2009	1
14	CAN-2012	1
15	CAN-2015	1
16	CAN-2018	1
40	FRA-2003	1
41	FRA-2006	1
42	FRA-2009	1
43	FRA-2012	1
44	FRA-2015	1
45	FRA-2018	1
46	DEU-2003	1
47	DEU-2006	1
48	DEU-2009	1
49	DEU-2012	1
50	DEU-2015	1
51	DEU-2018	1
70	IRL-2003	1
71	IRL-2006	1
72	IRL-2009	1
73	IRL-2012	1
105	NLD-2003	1
123	PRT-2003	1
124	PRT-2006	1
125	PRT-2009	1
126	PRT-2012	1

127	PRT-2015	1
158	USA-2003	1
159	USA-2006	1
160	USA-2009	1
161	USA-2012	1
162	USA-2015	1
163	USA-2018	1
164	GBR-2006	1
165	GBR-2009	1
166	GBR-2012	1
167	GBR-2015	1
168	GBR-2018	1

	country	cluster_label
23	DNK-2003	2
24	DNK-2006	2
25	DNK-2009	2
26	DNK-2012	2
27	DNK-2015	2
28	DNK-2018	2
37	FIN-2012	2
38	FIN-2015	2
39	FIN-2018	2
52	GRC-2003	2
53	GRC-2006	2
54	GRC-2009	2
55	GRC-2012	2
56	GRC-2015	2
57	GRC-2018	2
64	ISL-2003	2
65	ISL-2006	2
66	ISL-2009	2
67	ISL-2012	2
68	ISL-2015	2
69	ISL-2018	2
74	IRL-2015	2
75	IRL-2018	2
76	ISR-2006	2
77	ISR-2009	2
78	ISR-2012	2
79	ISR-2015	2
80	ISR-2018	2
81	ITA-2003	2
82	ITA-2006	2
83	ITA-2009	2
84	ITA-2012	2
85	ITA-2015	2
86	ITA-2018	2
101	LUX-2009	2
102	LUX-2012	2

103	LUX-2015	2
104	LUX-2018	2
106	NLD-2006	2
107	NLD-2009	2
108	NLD-2012	2
109	NLD-2015	2
110	NLD-2018	2
111	NOR-2003	2
112	NOR-2006	2
113	NOR-2009	2
114	NOR-2012	2
115	NOR-2015	2
116	NOR-2018	2
128	PRT-2018	2
133	SVK-2015	2
140	ESP-2003	2
141	ESP-2006	2
142	ESP-2009	2
143	ESP-2012	2
144	ESP-2015	2
145	ESP-2018	2
146	SWE-2003	2
147	SWE-2006	2
148	SWE-2009	2
149	SWE-2012	2
150	SWE-2015	2
151	SWE-2018	2

```
# perform hierarchical clustering
```

```
lbls = np.array(bestmodel_df.index)
```

```
cluster = linkage(bestmodel_df, method = 'complete', metric =  
"euclidean")
```

```
plt.figure(figsize=(18, 7))
```

```
dendrogram(cluster,  
            orientation='top',  
            labels=lbls,  
            distance_sort='descending')
```

```
plt.axhline(6,c="black",linewidth = 2)
```

```
plt.title("Euclidean Distance")
```

```
plt.show()
```

```
print()
```

```
# print out the clusters, using K=4
```


```
labels = fcluster(cluster, t=6.0, criterion='distance')
```

```
for k in np.arange(1,1+len(np.unique(labels))):
```

```
    print("group",k)
```

```
    print(bestmodel_df[labels==k].index.values)
```

```
    print('\n')
```



group 1

```
[ 'LVA-2003' 'LVA-2006' 'LVA-2009' 'USA-2003' 'USA-2006' 'USA-2009'
  'USA-2012' 'USA-2015' 'USA-2018' ]
```

group 2

```
[ 'GRC-2003' 'GRC-2006' 'GRC-2009' 'GRC-2012' 'GRC-2015' 'GRC-2018'
  'ISR-2006' 'ISR-2009' 'ISR-2012' 'ISR-2015' 'ISR-2018' 'ITA-2003'
  'ITA-2006' ]
```

group 3

```
[ 'AUS-2003' 'AUS-2006' 'AUS-2009' 'AUS-2012' 'AUS-2015' 'AUS-2018'
  'AUT-2003' 'AUT-2006' 'AUT-2012' 'AUT-2015' 'AUT-2018' 'CAN-2003'
  'CAN-2006' 'CAN-2009' 'CAN-2012' 'CAN-2015' 'CAN-2018' 'CZE-2003'
  'CZE-2006' 'CZE-2009' 'CZE-2012' 'CZE-2015' 'CZE-2018' 'DNK-2003'
  'DNK-2006' 'DNK-2009' 'DNK-2012' 'DNK-2015' 'DNK-2018' 'FIN-2003'
  'FIN-2006' 'FIN-2009' 'FIN-2012' 'FIN-2015' 'FIN-2018' 'FRA-2003'
  'FRA-2006' 'FRA-2009' 'FRA-2012' 'FRA-2015' 'FRA-2018' 'DEU-2003'
  'DEU-2006' 'DEU-2009' 'DEU-2012' 'DEU-2015' 'DEU-2018' 'HUN-2012'
  'HUN-2015' 'HUN-2018' 'ISL-2003' 'ISL-2006' 'ISL-2009' 'ISL-2012'
  'ISL-2015' 'ISL-2018' 'IRL-2003' 'IRL-2006' 'IRL-2009' 'IRL-2012'
  'IRL-2015' 'IRL-2018' 'ITA-2009' 'ITA-2012' 'ITA-2015' 'ITA-2018'
  'LUX-2003' 'LUX-2006' 'LUX-2009' 'LUX-2012' 'LUX-2015' 'LUX-2018'
  'NLD-2003' 'NLD-2006' 'NLD-2009' 'NLD-2012' 'NLD-2015' 'NLD-2018'
  'NOR-2003' 'NOR-2006' 'NOR-2009' 'NOR-2012' 'NOR-2015' 'NOR-2018'
  'POL-2003' 'POL-2006' 'POL-2009' 'POL-2012' 'POL-2015' 'POL-2018'
  'PRT-2003' 'PRT-2006' 'PRT-2009' 'PRT-2012' 'PRT-2015' 'PRT-2018'
  'SVN-2006' 'SVN-2009' 'SVN-2012' 'SVN-2015' 'SVN-2018' 'ESP-2003'
  'ESP-2006' 'ESP-2009' 'ESP-2012' 'ESP-2015' 'ESP-2018' 'SWE-2003'
  'SWE-2006' 'SWE-2009' 'SWE-2012' 'SWE-2015' 'SWE-2018' 'CHE-2003'
  'CHE-2006' 'CHE-2009' 'CHE-2012' 'CHE-2015' 'CHE-2018' 'GBR-2006'
  'GBR-2009' 'GBR-2012' 'GBR-2015' 'GBR-2018' ]
```

group 4

```
[ 'EST-2006' 'EST-2009' 'EST-2012' 'EST-2015' 'EST-2018' 'HUN-2003'  
  'HUN-2006' 'HUN-2009' 'KOR-2003' 'KOR-2006' 'KOR-2009' 'KOR-2012'  
  'KOR-2015' 'KOR-2018' 'LVA-2012' 'LVA-2015' 'LVA-2018' 'SVK-2003'  
  'SVK-2006' 'SVK-2009' 'SVK-2012' 'SVK-2015' 'SVK-2018' ]
```
